

A Presentation on *A Performance Survey on Stack-based and Register-based Virtual Machines*

Alex Fang, Tyler Liu

LA SALLE COLLEGE HIGH SCHOOL

November 30, 2016

# Problem Statement

## Fact

*Every programmer deals with interpreted languages (Java, Python, etc.) every day.*

1. Arguments arise when discussing the performance of interpreted languages: WHICH is faster? WHY and HOW is it faster?
2. As programmers we want to provide a better understanding of this issue using a SCIENTIFIC METHOD

# Problem Boil-Down

## Definition

The performance of an interpreted programming language = The runtime performance of its virtual machine

## Fact

*Two most popular types of virtual machines: The Stack-based Virtual Machine and The Register-based Virtual Machine*

Through comparing and contrasting the performance between a stack-based virtual machine and a register-based virtual machine, we can find out:

1. Which one has the best performance
2. Different best-case applications of the two architectures

# Hypothesis

Formula on the theoretical performance differences between a stack-based machine and a register based machine, proposed by Davis et. al:

$$T_{VRM} \approx T_{VSM} - \#dispatches \times T_{dispatch} + \#fetches \times T_{fetch}$$

- ✚ *VRM*: Register-based Virtual Machine
- ✚ *VSM*: Stack-based Virtual Machine
- ✚ *T*: Time
- ✚ *#dispatches*: Total time spent in dispatches
- ✚ *#fetches*: Total time spent in fetches

# Reasoning Behind The Hypothesis

$$T_{VRM} \approx T_{VSM} - \#dispatches \times T_{dispatch} + \#fetches \times T_{fetch}$$

Stack-based Bytecode requires only ONE operand per instruction, whereas...

Register-based bytecode requires THREE operands per instruction.

**Result:** The stack-based virtual machine has a performance advantage in FETCHING the operands (fetch time); the register-based virtual machine has a advantage in a higher instruction efficiency (DISPATCHING instructions in general)

## ✚ Estimates:

- ◇ Stack VM can perform better in fetch performance since it has 1-2 less operands per instruction
- ◇ Register VM can have better performance in dispatch amount since it concentrates amount of instructions by having explicitly specified memory addresses

# Approach

**Instrument.** We constructed two new virtual machines from scratch (one stack-based, another register-based)

**Benchmark.** We demonstrate the performance of each machines through two different types of benchmarks:

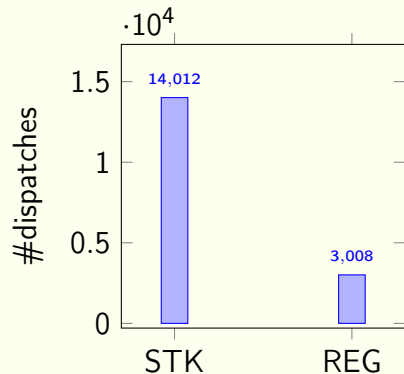
- ✦ Regular algorithm-based benchmark: Fibonacci, Recursion
- ✦ High density arithmetic benchmark: ExhaustiveCollatz, AddictiveAddition

**Measurements.** Using the formula of Davis et. al, we measured the four core factors of performance:


- ✦ Amount of dispatches per program
- ✦ Total time spent in instruction dispatch
- ✦ Total time spent in fetching instruction operands
- ✦ Total execution time

# Benchmark Data - Regular Algorithm-based Benchmark

Fibonacci - Amount of Dispatches



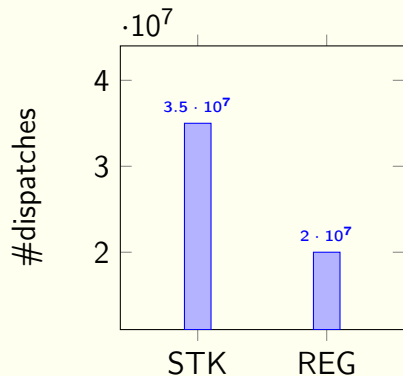
Fibonacci - Execution Time



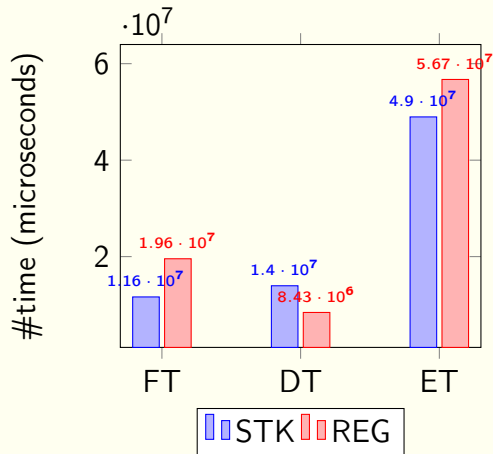
Note: STK=Stack-based Virtual Machine, REG=Register-based Virtual Machine

# Benchmark Data - High Density Arithmetic Benchmark

AddictiveAddition - Amount of Dispatches



AddictiveAddition - Execution Time



Note: STK=Stack-based Virtual Machine, REG=Register-based Virtual Machine



# Conclusion

- ✦ Overall, a register-based virtual machine is around 20% faster
- ✦ Stack-based VM performed better in fetch time (less)
- ✦ If you want to implement a high-performance, compact Domain-Specific Language on limited hardware, go for a register-based virtual machine!
- ✦ If you favor simplicity (both in byte code and in code for VM) over performance and want to perform dense read/write to the VM's memory space, implement a stack-based virtual machine!

## Extra Works Conducted

- ✦ We pre-published our paper on the academia's largest e-print pre-publishing website, arXiv.org.
- ✦ We opensourced Everything related to this paper: virtual machine programs, bytecodes... In the hope of a better transparency and more feedback
- ✦ We plan to publish a revised paper on a Computer Science-related transaction after this conference
- ✦ Future works are planned for us to make revisions to our method of analysis to make the data gained more consistent and scientific.

- ✚ Brian Davis, Andrew Beatty, Kevin Casey, David Gregg, and John Waldron. The case for virtual register machines. In *Proceedings of the 2003 Workshop on Interpreters, Virtual Machines, and Emulators*, IVME '03. pages 41-49, New York, NY, USA, 2003. ACM.